

# Tidyverse introduction

Akitaka Matsuo

# Introduction

- Data manipulation (wrangling) takes the vast majority of time in data science
  - That is more true when the data are big
- **tidyverse** provides a unified method for data wrangling
  - “The tidyverse is an **opinionated collection of R packages** designed for data science. All packages share an underlying design philosophy, grammar, and data structures.”



Harvard  
Business  
Review

DATA

## Data Scientist: The Sexiest Job of the 21st Century

by [Thomas H. Davenport](#) and [D.J. Patil](#)

From the October 2012 Issue

Summary

Save

Share

<sup>17</sup> Comment

Print

**\$8.95** Buy Copies



When Jonathan Goldman arrived for work in June 2006 at LinkedIn, the business networking site, the

# Why or why not **tidyverse**?

## Pros:

- Once getting used to it, you can do a lot of complicated data management work efficiently
- Improved reproducibility of research
  - Easy to read data manipulation process
- Unified syntax to access to many background options
  - **dbplyr**: SQL databases
  - **dtplyr**: data.table
  - **sparklyr**: Apache Spark (Day 5)

## Cons:

- Opinionated
- Steep learning curve

# tidyverse packages

**tidyverse** (tidy + universe) is a group of packages working together.

- **dplyr**: “a grammar of data manipulation”
- **tidyr**: data reshaping to make the data *tidy*
- **readr**: read various data sets (also **readxl**)
- **purrr**: `map()`
- **ggplot2**: data visualization

We will mostly look at **dplyr** and **tidyr**.

# Load it

```
library(tidyverse)
## -- Attaching packages ----- tidyverse 1.3.0 --
## v ggplot2 3.3.2      v purrr 0.3.4
## v tibble 3.0.2       v dplyr 1.0.0
## v tidyr 1.1.0        v stringr 1.4.0
## v readr 1.3.1        v forcats 0.5.0
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

# dplyr

A package to provide a new grammar to carry out data manipulation

With `dplyr` we don't have to...

```
mtcars_sub <- subset(mtcars, cyl == 4)
```

```
mtcars_sub$log_hp <- log(mtcars_sub$hp)
```

```
model_regression <- lm(mpg ~ log_hp + wt, data =  
mtcars_sub)
```

```
summary(model_regression)
```

# The grammar of `dplyr`

## Basics:

- `%>%`: pipe operator
- `filter()`: select rows
- `select()`: select variables
- `mutate()`: manipulate variables

## Group based operators:

- `group_by()`: group the dataset
- `summarise()`: get a summary value for the group
- `count()`: table-like function

# %>% - pipe operator

- %>% sends the results from one function to the another. For example

```
x <- 1:10  
x %>% sum()
```

- The pipe command, send `x` to `sum()` as the first argument of the function (i.e. `sum(x)`)

- When you want to send to the second or later argument of a function, you can use `.` indicator

```
mtcars %>% lm(mpg ~ wt + cyl + hp, data = .)
```

- You can use multiple %>% to create a chain of operations

```
mtcars %>% lm(mpg ~ wt + cyl + hp, data = .) %>%  
summary()
```

- The pipe operator is one of the main reasons why tidyverse codes look different from base-R



# %>% - pipe outputs

```
mtcars %>% lm(mpg ~ wt + cyl + hp, data = .) %>% summary()
##
## Call:
## lm(formula = mpg ~ wt + cyl + hp, data = .)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.9290 -1.5598 -0.5311  1.1850  5.8986
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  38.75179     1.78686   21.687  < 2e-16 ***
## wt          -3.16697     0.74058   -4.276 0.000199 ***
## cyl         -0.94162     0.55092   -1.709 0.098480 .
## hp          -0.01804     0.01188   -1.519 0.140015
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.512 on 28 degrees of freedom
## Multiple R-squared:  0.8431, Adjusted R-squared:  0.8263
## F-statistic: 50.17 on 3 and 28 DF,  p-value: 2.184e-11
```

# filter()

- `filter()` subset the data, using the conditional statement.

```
filter(mtcars, cyl == 4)
```

- Or, using pipe operator

```
mtcars %>% filter(cyl == 4)
```

# select()

`select()` is used for selecting variables to keep or drop from a data.frame.

## Selecting

```
mtcars <- mtcars %>% rownames_to_column(var = "car_name")
```

```
mtcars_sub <- mtcars %>%  
  select(car_name, mpg, cyl, wt, hp)
```

```
mtcars_sub %>% head()  
##           car_name  mpg  cyl    wt  hp  
## 1      Mazda RX4  21.0    6 2.620 110  
## 2   Mazda RX4 Wag  21.0    6 2.875 110  
## 3    Datsun 710  22.8    4 2.320  93  
## 4  Hornet 4 Drive  21.4    6 3.215 110  
## 5 Hornet Sportabout 18.7    8 3.440 175  
## 6      Valiant  18.1    6 3.460 105
```

# select()

## Dropping

```
mtcars_sub2 <- mtcars %>%  
  select(-c(gear, carb, disp, qsec))
```

```
mtcars_sub2 %>% head()
```

##		car_name	mpg	cyl	hp	drat	wt	vs	am
## 1		Mazda RX4	21.0	6	110	3.90	2.620	0	1
## 2		Mazda RX4 Wag	21.0	6	110	3.90	2.875	0	1
## 3		Datsun 710	22.8	4	93	3.85	2.320	1	1
## 4		Hornet 4 Drive	21.4	6	110	3.08	3.215	1	0
## 5		Hornet Sportabout	18.7	8	175	3.15	3.440	0	0
## 6		Valiant	18.1	6	105	2.76	3.460	1	0

# mutate()

`mutate()`: manipulation (either generation or alteration of variable). Let's say you want to create a dummy variable to indicate a car with above average mpg.

```
mtcars <- mtcars %>%  
  mutate(good_mpg = (mpg > mean(mpg)))  
  #or mutate(good_mpg = (mpg > mean(mpg)) %>% as.integer())
```

```
mtcars %>% select(car_name, mpg, good_mpg) %>% head()  
##           car_name  mpg good_mpg  
## 1      Mazda RX4  21.0      TRUE  
## 2    Mazda RX4 Wag  21.0      TRUE  
## 3    Datsun 710  22.8      TRUE  
## 4   Hornet 4 Drive  21.4      TRUE  
## 5 Hornet Sportabout 18.7     FALSE  
## 6      Valiant  18.1     FALSE
```

# arrange, slice

```
mtcars %>% arrange(desc(mpg)) %>%  
  select(car_name, mpg, good_mpg) %>%  
  slice(1:5)
```

```
##           car_name  mpg good_mpg  
## 1 Toyota Corolla 33.9      TRUE  
## 2      Fiat 128 32.4      TRUE  
## 3   Honda Civic 30.4      TRUE  
## 4 Lotus Europa 30.4      TRUE  
## 5    Fiat X1-9 27.3      TRUE
```

# group\_by, summarise

`dplyr` also provides a very powerful functionality to get an aggregate measure for specific groups of the interest. For instance suppose that you want to know an average `mpg` by number of `cyl`. You can achieve that with the combination of `group_by()` and `summarise()`

```
mtcars %>%  
  group_by(cyl) %>% # this can be more than one grouping variable  
  summarise(avg_mpg = mean(mpg, na.rm = T),  
            count = n()) # and you can generate more than one  
variable  
## `summarise()` ungrouping output (override with `.groups` argument)  
## # A tibble: 3 x 3  
##   cyl avg_mpg count  
##   <dbl> <dbl> <int>  
## 1     4    26.7     11  
## 2     6    19.7      7  
## 3     8    15.1     14
```

# group\_by, mutate

- `group_by()` can work with `mutate()` and others

```
mtcars %>%  
  group_by(cyl) %>%  
  mutate(good_mpg = (mpg > mean(mpg))) %>%  
  arrange(desc(mpg)) %>%  
  slice(1) %>%  
  select(car_name, mpg, good_mpg)  
## Adding missing grouping variables: `cyl`  
## # A tibble: 3 x 4  
## # Groups:   cyl [3]  
##       cyl car_name      mpg good_mpg  
##   <dbl> <chr>      <dbl> <lgl>  
## 1     4 Toyota Corolla   33.9 TRUE  
## 2     6 Hornet 4 Drive   21.4 TRUE  
## 3     8 Pontiac Firebird  19.2 TRUE
```



# What is tidy data?

Essentially speaking, the tidy data are

- Columns are variables
- Rows are observations

# What is tidy data?

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	216766	1280425583

variables

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	216766	1280425583

observations

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	216766	1280425583

values

# What is NOT tidy data?

These are historical household data in the US States.

```
data_income <-  
  read_csv("data/data_state_  
income.csv")  
data_income %>%  
  select(1:5) %>%  
    slice(1:6) %>%  
    knitr::kable()
```

How to make it tidy?

state	Y2018	Y2017	Y2016	Y2015
Alabama	49936	50865	47221	44509
Alaska	68734	77987	75723	75112
Arizona	62283	59700	57100	52248
Arkansas	49781	49751	45907	42798
California	70489	70038	66637	63636
Colorado	73034	74984	70566	66596

# The data in tidy format

This is the same data in tidy format.  
For that kind of reshaping between  
wide to long format, we can use  
`tidyr::pivot_longer`

state	year	income
Alabama	2018	49936
Alabama	2017	50865
Alabama	2016	47221
Alabama	2015	44509
Alabama	2014	42278
Alabama	2013	47320

# The data in tidy format

- This is an code example to generate the previous pages output

```
data_income_long <- data_income %>%  
  pivot_longer(Y2018:Y1984, names_to = "year",  
               values_to = "income") %>%  
  mutate(year = str_sub(year, 2) %>% as.integer())
```

```
data_income_long %>%  
  slice(1:3)  
## # A tibble: 3 x 3  
##   state      year income  
##   <chr>    <int>  <dbl>  
## 1 Alabama  2018   49936  
## 2 Alabama  2017   50865  
## 3 Alabama  2016   47221
```

# Going back to the wide format

```
data_income_long %>%  
  pivot_wider(names_from = year, values_from = income) %>%  
  slice(1:10) %>% select(1:8)
```

## # A tibble: 10 x 8

	state	`2018`	`2017`	`2016`	`2015`	`2014`	`2013`	`2012`
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	Alabama	49936	50865	47221	44509	42278	47320	43464
## 2	Alaska	68734	77987	75723	75112	67629	72472	63648
## 3	Arizona	62283	59700	57100	52248	49254	52611	47044
## 4	Arkansas	49781	49751	45907	42798	44922	39376	39018
## 5	California	70489	70038	66637	63636	60487	60794	57020
## 6	Colorado	73034	74984	70566	66596	60940	67912	57255
## 7	Connecticut	72812	74304	75923	72889	70161	69291	64247
## 8	Delaware	65012	64961	58046	57756	57522	54091	48972
## 9	D.C.	85750	81282	70982	70071	68277	60057	65246
## 10	Florida	54644	53086	51176	48825	46140	48532	46071

# Summary

## What we've learned:

- What is `tidyverse`
- Why we should learn
- `dplyr`
  - `filter()`, `select()`, `mutate()`
  - `group_by()`, `summarise()`
- `tidyr`
  - `pivot_wider()`

Next: `data.table`, reading data efficiently